



## LINEAR CLASSIFIERS IN PYTHON

# Support Vectors


Michael (Mike) Gelbart

Instructor

The University of British Columbia

# What is an SVM?

- Linear classifiers (so far)
- Trained using the hinge loss and L2 regularization






# What are support vectors?

- Support vector: a training example **not** in the flat part of the loss diagram
- Support vector: an example that is incorrectly classified **or** close to the boundary
- If an example is not a support vector, removing it has no effect on the model
- Having a small number of support vectors makes kernel SVMs really fast




# Max-margin viewpoint

- The SVM maximizes the "margin" for linearly separable datasets
- Margin: distance from the boundary to the closest points



# Max-margin viewpoint

- The SVM maximizes the "margin" for linearly separable datasets
- Margin: distance from the boundary to the closest points





## LINEAR CLASSIFIERS IN PYTHON

**Let's practice!**



## LINEAR CLASSIFIERS IN PYTHON


# Kernel SVMs

Michael (Mike) Gelbart

Instructor


The University of British Columbia

# Transforming your features

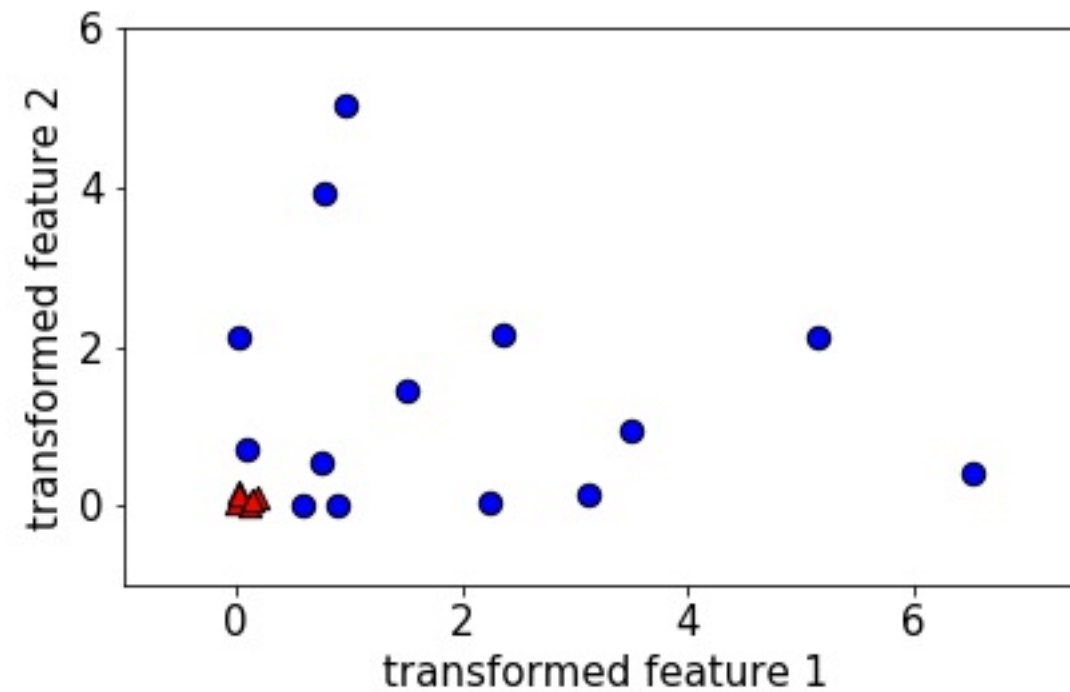





# Transforming your features

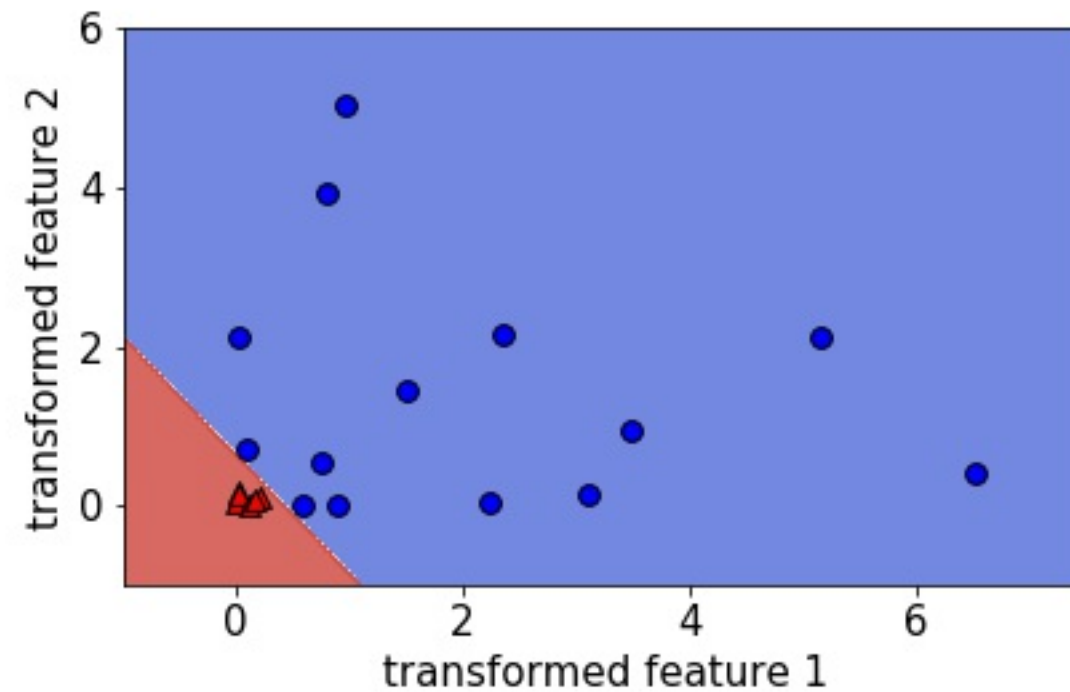



# Transforming your features




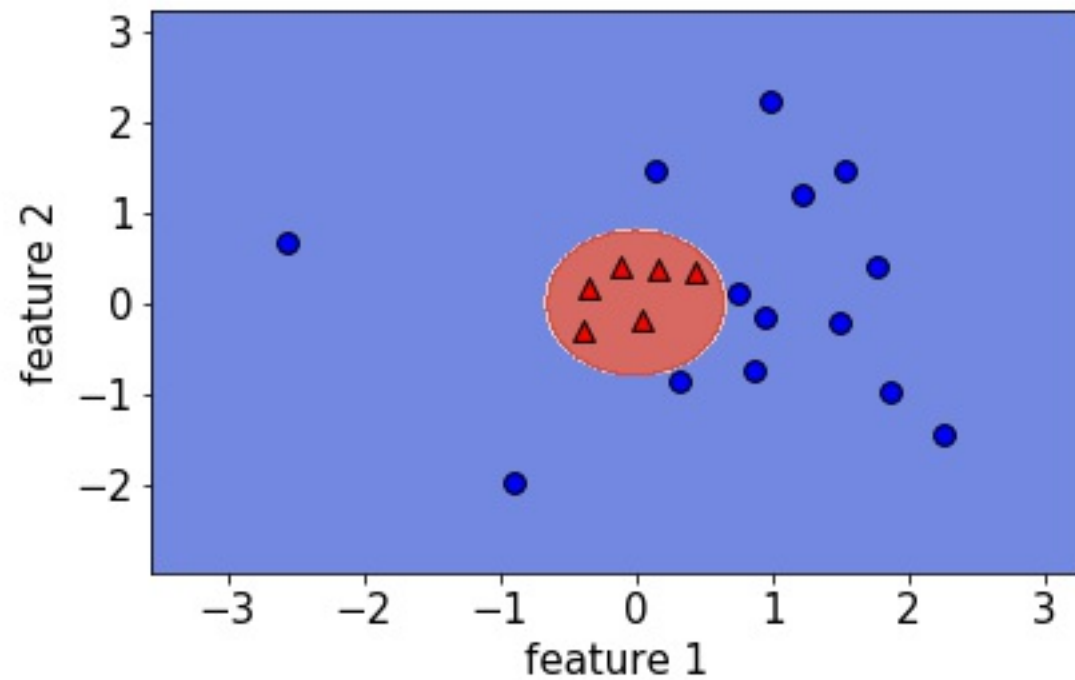
$$\text{transformed feature} = (\text{original feature})^2$$

# Transforming your features



$$\text{transformed feature} = (\text{original feature})^2$$

# Transforming your features




$$\text{transformed feature} = (\text{original feature})^2$$

# Kernel SVMs

```
In [1]: from sklearn.svm import SVC
```


```
In [2]: svm = SVC(gamma=1)      # default is kernel="rbf"
```



# Kernel SVMs

```
In [1]: from sklearn.svm import SVC
```

```
In [2]: svm = SVC(gamma=0.01) # default is kernel="rbf"
```




- smaller gamma leads to smoother boundaries

# Kernel SVMs

```
In [1]: from sklearn.svm import SVC
```

```
In [2]: svm = SVC(gamma=2)      # default is kernel="rbf"
```



- larger gamma leads to more complex boundaries



## LINEAR CLASSIFIERS IN PYTHON

**Let's practice!**





LINEAR CLASSIFIERS IN PYTHON

# Comparing logistic regression and SVM

Michael (Mike) Gelbart

Instructor

The University of British Columbia



# Pros and Cons

## Logistic regression:

- Is a linear classifier
- Can use with kernels, but slow
- Outputs meaningful probabilities
- Can be extended to multi-class
- All data points affect fit
- L2 or L1 regularization

## Support vector machine (SVM):

- Is a linear classifier
- Can use with kernels, and fast
- Does not naturally output probabilities
- Can be extended to multi-class
- Only "support vectors" affect fit
- Conventionally just L2 regularization

# Use in scikit-learn

Logistic regression in sklearn:

- `linear_model.LogisticRegression`

Key hyperparameters in sklearn:

- `C` (inverse regularization strength)
- `penalty` (type of regularization)
- `multi_class` (type of multi-class)

SVM in sklearn:

- `svm.LinearSVC` and `svm.SVC`

Key hyperparameters in sklearn:

- `C` (inverse regularization strength)
- `kernel` (type of kernel)
- `gamma` (inverse RBF smoothness)



# SGDClassifier

- SGDClassifier: scales well to large datasets

```
In [1]: from sklearn.linear_model import SGDClassifier
```

```
In [2]: logreg = SGDClassifier(loss='log')
```

```
In [3]: linsvm = SGDClassifier(loss='hinge')
```

- SGDClassifier hyperparameter alpha is like  $1/C$



## LINEAR CLASSIFIERS IN PYTHON

**Let's practice!**



## LINEAR CLASSIFIERS IN PYTHON

# Conclusion

Michael (Mike) Gelbart

Instructor

The University of British Columbia



# How does this course fit into Data Science?

- Data science
  - --> Machine learning
    - --> --> Supervised learning
      - --> --> --> Classification
        - --> --> --> --> Linear classifiers (this course)



## LINEAR CLASSIFIERS IN PYTHON

**Congratulations &  
Thanks!**